

Midterm Review

Midterm Exam

- ❖ Time: see syllabus in Canvas
- ❖ Location: Canvas + Honorlock
 - <https://fiuhelp.force.com/canvas/s/article/Honorlock-students>
- ❖ Scope
 - Chapter 1 – Introduction
 - Chapter 2 – Application layer
 - Chapter 3 – Transport layer

Midterm Exam

❖ Question format


- 20 questions (4 points each), similar to quiz questions
- 4 problems (5 points each), similar to homework problems

❖ How to prepare

- Review slides
- Quiz questions
- Homework problems

Chapter I

Introduction



*Computer
Networking: A Top
Down Approach*

Slides adopted from original ones provided by the textbook authors.

Chapter 1: roadmap

1.1 *what is the Internet?*

1.2 network edge

- end systems, access networks, links

1.3 network core

- packet switching, circuit switching, network structure


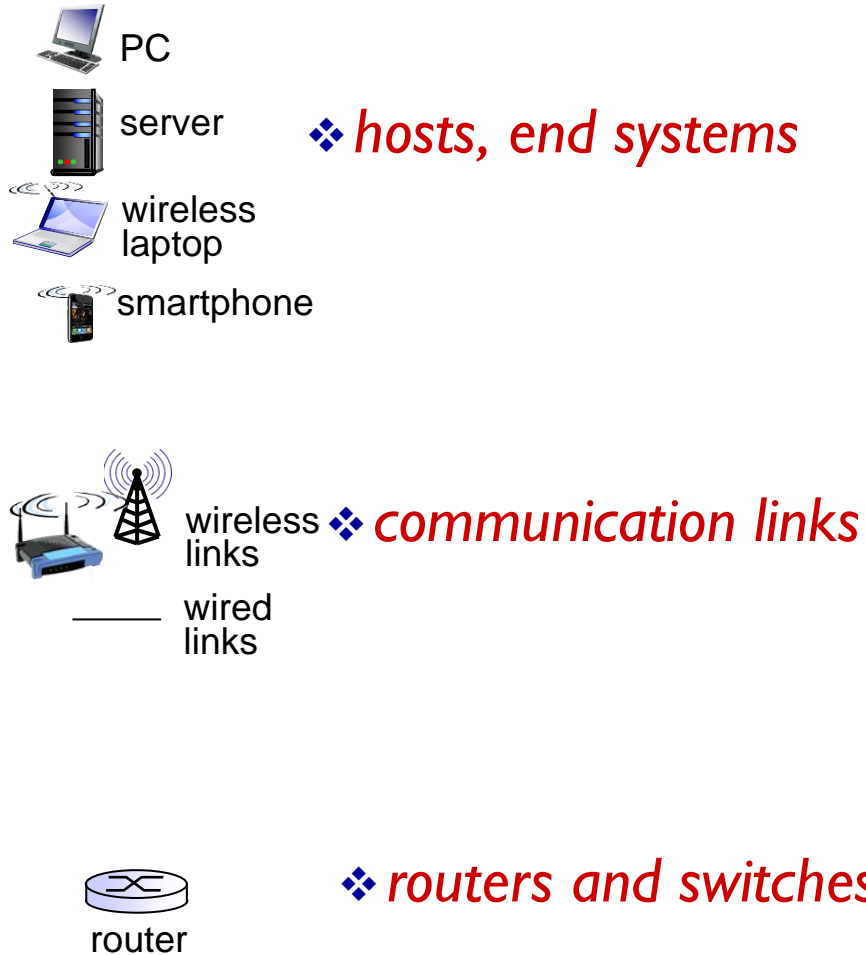
1.4 delay, loss, throughput in networks

1.5 protocol layers, service models

1.6 networks under attack: security

1.7 history

What's the Internet



Chapter 1: roadmap

1.1 what is the Internet?

1.2 network edge

- end systems, access networks, links

1.3 network core

- packet switching, circuit switching, network structure

1.4 delay, loss, throughput in networks

1.5 protocol layers, service models


1.6 networks under attack: security

1.7 history

A closer look at network structure:

❖ *network edge:*

- hosts: clients and servers
- servers often in data centers



Access networks

- ❖ DSL: several Mbps, dedicated access
- ❖ Cable: tens of Mbps, shared access
- ❖ Ethernet: Gbps, for institutional networks
- ❖ Wireless: WIFI, 3G/4G cellular

Physical Media

❖ guided media

- Twisted pair: Ethernet
- Coax: cable networks
- Fiber: optical networks

❖ unguided media

- terrestrial microwave
- LAN (e.g., Wifi)
- wide-area (e.g., cellular)
- satellite

Chapter 1: roadmap

1.1 what is the Internet?

1.2 network edge

- end systems, access networks, links

1.3 network core

- circuit switching, packet switching, network structure

1.4 delay, loss, throughput in networks


1.5 protocol layers, service models

1.6 networks under attack: security

1.7 history

The network core


- ❖ mesh of interconnected routers
- ❖ the fundamental question:
how is data transferred through net?
 - **circuit switching**: dedicated circuit per call: telephone net, GSM
 - **packet-switching**: data sent thru net in discrete “chunks”



Circuit switching

end-end resources allocated to, reserved for “call” between source & dest:

- ❖ dedicated resources: no sharing
- ❖ circuit-like (guaranteed) performance
- ❖ call setup required



Network Core: Packet Switching

each end-end data stream
divided into *packets*

- ❖ user A, B packets *share* network resources
- ❖ each packet uses full link bandwidth
- ❖ resources used *as needed*

resource contention:

- ❖ aggregate resource demand can exceed amount available
- ❖ congestion: packets queue, wait for link use
- ❖ store and forward: packets move one hop at a time

Packet switching versus circuit switching

	Circuit switching	Packet switching
Pros	<ul style="list-style-type: none">• Performance guarantees	<ul style="list-style-type: none">• High resource utilization• No call set up
Cons	<ul style="list-style-type: none">• Low resource utilization• Call set up	<ul style="list-style-type: none">• Congestion

Chapter 1: roadmap

1.1 what *is* the Internet?

1.2 network edge

- end systems, access networks, links

1.3 network core

- packet switching, circuit switching, network structure


1.4 delay, loss, throughput in networks

1.5 protocol layers, service models

1.6 networks under attack: security

1.7 history

Four sources of packet delay



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

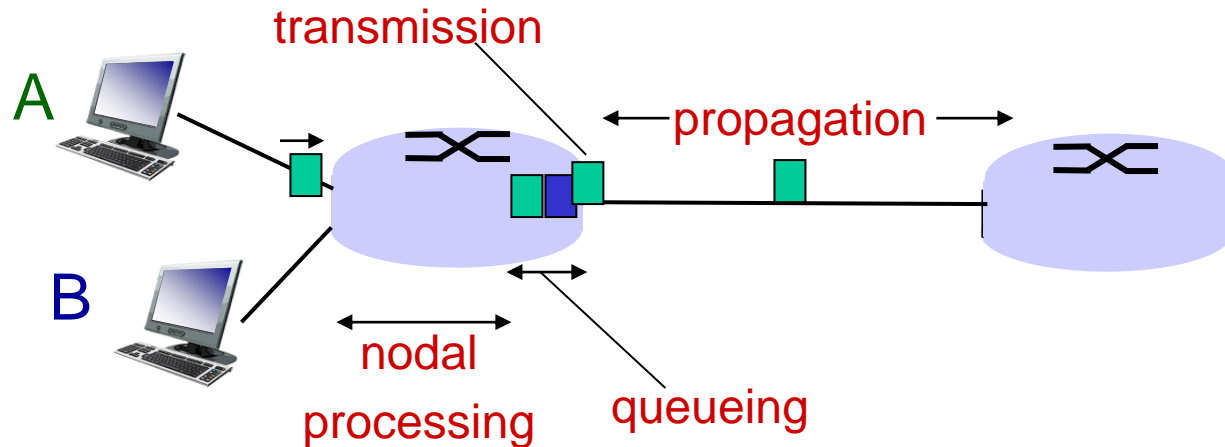
d_{proc} : nodal processing

- check bit errors
- determine output link
- typically < msec

d_{queue} : queueing delay

- time waiting at output link for transmission
- depends on congestion level of router

Four sources of packet delay



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

d_{trans} : transmission delay:

- L : packet length (bits)
- R : link bandwidth (bps)
- $d_{\text{trans}} = L/R$

d_{prop} : propagation delay:

- d : length of physical link
- s : propagation speed in medium ($\sim 2 \times 10^8$ m/sec)
- $d_{\text{prop}} = d/s$

Throughput

- ❖ *throughput*: rate (bits/time unit) at which bits transferred between sender/receiver
 - *instantaneous*: rate at given point in time
 - *average*: rate over longer period of time
- ❖ determined by **bottleneck link**
 - link on end-end path that constrains end-end throughput

Chapter 1: roadmap

1.1 what is the Internet?

1.2 network edge

- end systems, access networks, links

1.3 network core

- packet switching, circuit switching, network structure

1.4 delay, loss, throughput in networks


1.5 protocol layers, service models

1.6 networks under attack: security

1.7 history


Internet protocol stack

- ❖ *application*: supporting network applications
 - FTP, SMTP, HTTP
- ❖ *transport*: process-process data transfer
 - TCP, UDP
- ❖ *network*: routing of datagrams from source to destination
 - IP, routing protocols
- ❖ *link*: data transfer between neighboring network elements
 - Ethernet, 802.111 (WiFi), PPP
- ❖ *physical*: bits “on the wire”



Chapter 2

Application Layer



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

Slides adopted from original ones provided by the textbook authors.

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Application architectures

❖ Client-server

- Always-on server, intermittently connected client.
- Servers are bottlenecks.

❖ Peer-to-peer (P2P)

- Peers intermittently connected.
- Highly scalable but difficult to manage.

Sockets

- ❖ Process sends/receives messages to/from its socket.
 - Processes are identified by IP addresses and TCP/UDP port numbers.
- ❖ OS provides APIs for creating sockets.

Internet transport protocols services

TCP service:

- ❖ *reliable transport* between sending and receiving process
- ❖ *flow control*: sender won't overwhelm receiver
- ❖ *congestion control*: throttle sender when network overloaded
- ❖ *does not provide*: timing, minimum throughput guarantee, security
- ❖ *connection-oriented*: setup required between client and server processes

UDP service:

- ❖ *unreliable data transfer* between sending and receiving process
- ❖ *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

HTTP overview

uses TCP:

- ❖ client initiates TCP connection (creates socket) to server, port 80
- ❖ server accepts TCP connection from client
- ❖ HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- ❖ TCP connection closed

HTTP is “stateless”

- ❖ server maintains no information about past client requests

HTTP connections


non-persistent HTTP

- ❖ at most one object sent over TCP connection
 - connection then closed
- ❖ downloading multiple objects required multiple connections
 - use parallel TCP connections to accelerate


persistent HTTP

- ❖ multiple objects can be sent over single TCP connection between client, server

HTTP request message: general format



HTTP response message: general format



User-server state: cookies

- ❖ Cookies help web sites remember use states.
- ❖ Four components:
 - 1) set-cookie header line in HTTP *response* message
 - 2) cookie header line in HTTP *request* message
 - 3) cookie file kept on user's host, managed by user's browser
 - 4) back-end database at Web site

Web caches

- ❖ **Goal:** satisfy client request without involving origin server
 - reduce response time and traffic
- ❖ user sets browser: Web accesses via cache, browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client
- ❖ Conditional GET: don't send object if cache has up-to-date cached version

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Electronic Mail

- ❖ SMTP
 - mail transfer protocol
 - client/server model
 - based on TCP, port 25
- ❖ Mail message format defined in RFC 822
 - header and body
- ❖ POP3 and IMAP
 - mail access protocols

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks


2.7 socket programming with UDP and TCP

DNS: Domain Name System

❖ DNS services

- hostname to IP address translation
- host aliasing: canonical, alias names
- mail server aliasing
- load distribution

Distributed, Hierarchical Database




- ❖ Root name servers
- ❖ Top-level domain (TLD) servers
- ❖ Authoritative DNS servers
- ❖ Local name server

DNS name resolution example

- ❖ host at cis.poly.edu wants IP address for gaia.cs.umass.edu

iterated query:

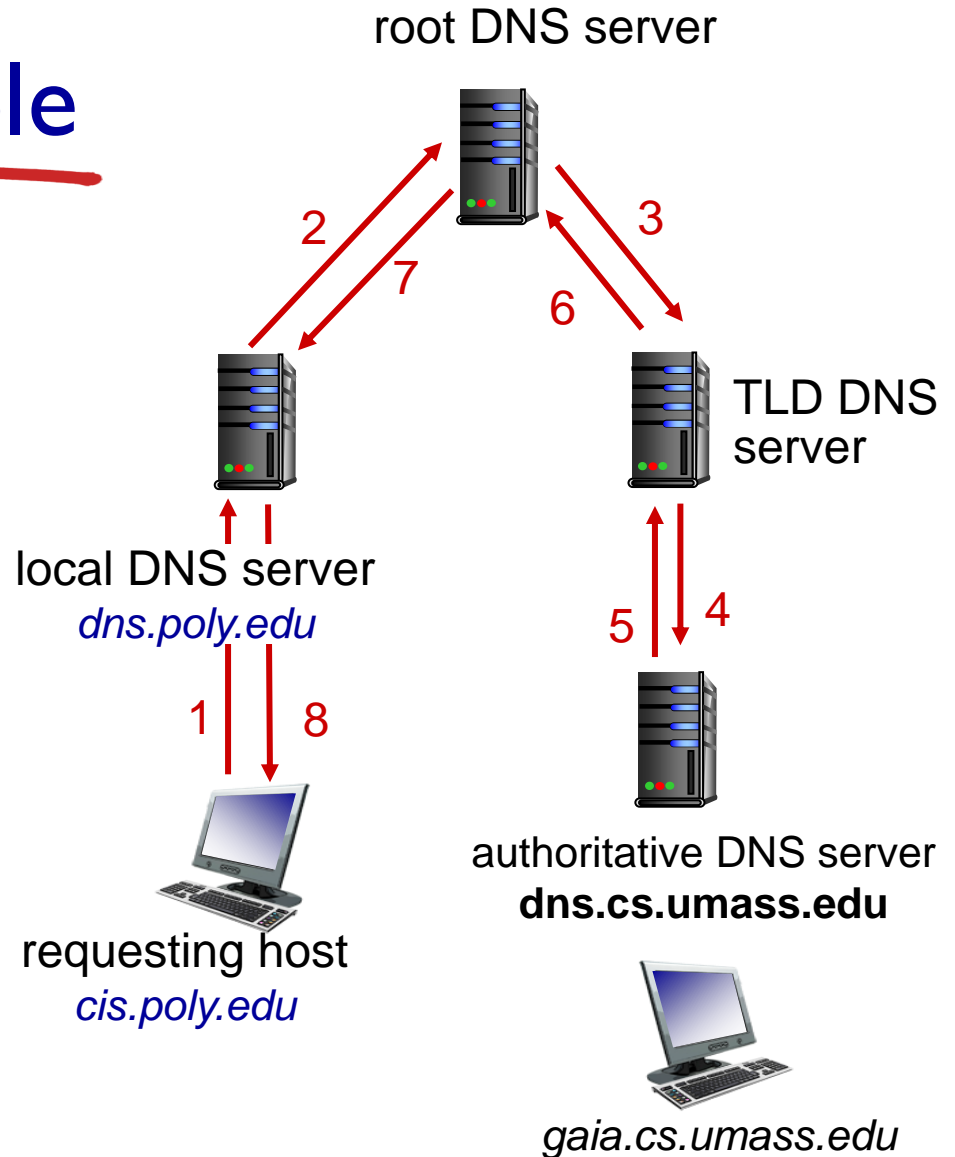
- ❖ contacted server replies with name of server to contact
- ❖ “I don’t know this name, but ask this server”



DNS name resolution example

recursive query:

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy



DNS: caching, updating records

- ❖ once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
- ❖ cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire

DNS records

DNS: distributed db storing resource records (RR)

RR format: (name, value, type, ttl)

type=A

- **name** is hostname
- **value** is IP address

type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

type=CNAME

- **name** is alias name for some “canonical” (the real) name
- **www.ibm.com** is really **servereast.backup2.ibm.com**
- **value** is canonical name

type=MX

- **value** is name of mailserver associated with **name**

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

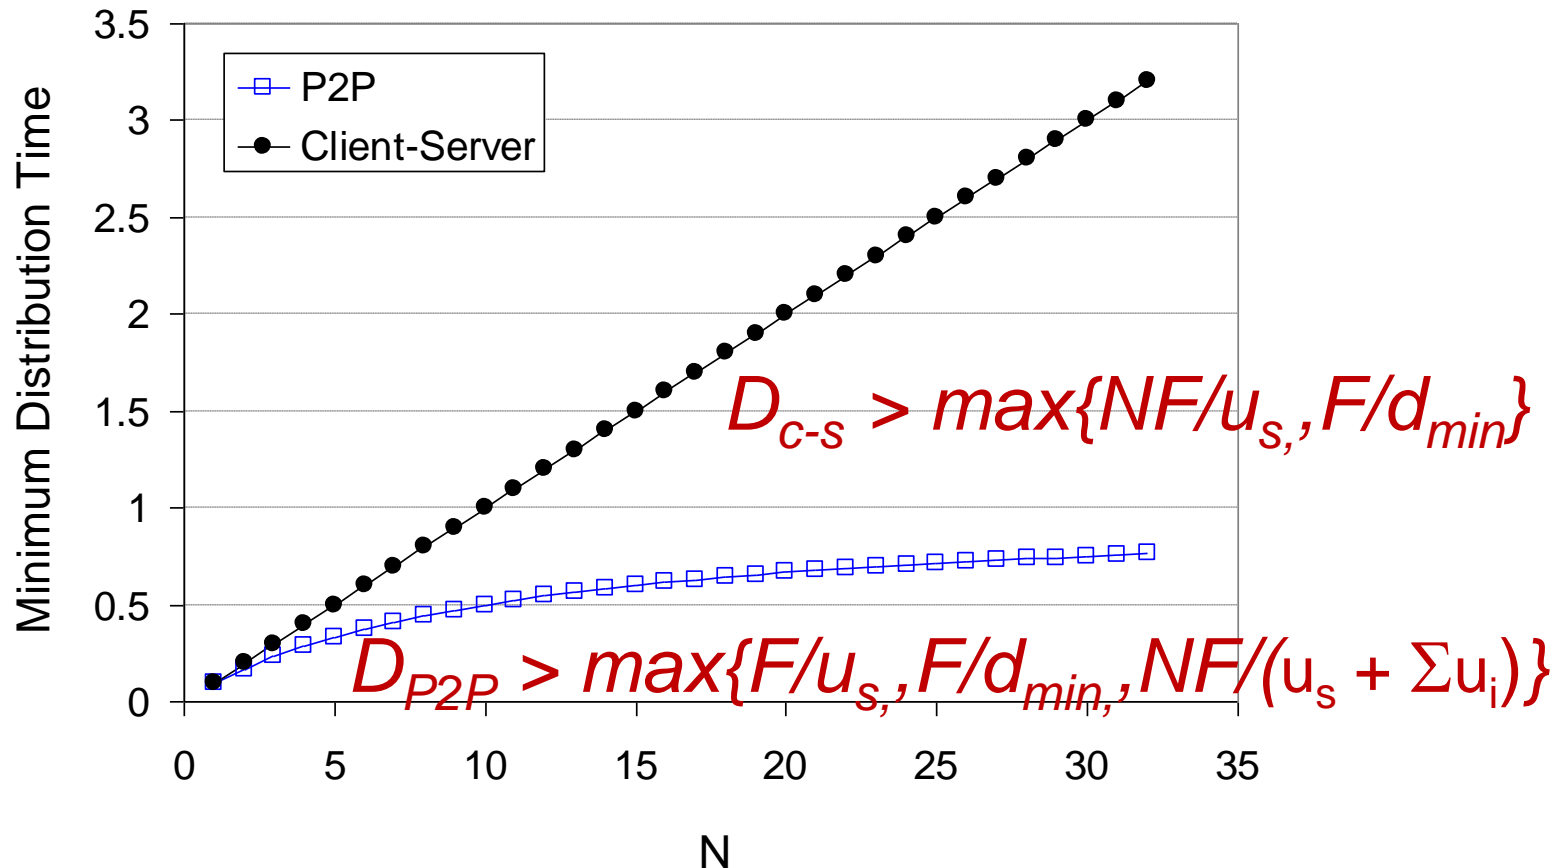
2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$




P2P file distribution: BitTorrent

- ❖ file divided into 256Kb chunks
- ❖ peers in torrent send/receive file chunks

tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file



BitTorrent: requesting, sending file chunks

requesting chunks:

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, Alice asks each peer for list of chunks that they have
- ❖ Alice requests missing chunks from peers, **rarest first**

*sending chunks: **tit-for-tat***

- ❖ Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- ❖ every 30 secs: randomly select another peer, starts sending chunks
 - “**optimistically unchoke**” this peer
 - newly chosen peer may join top 4

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks (CDNs)

2.7 socket programming with UDP and TCP

Streaming multimedia: DASH

❖ *DASH*: *D*ynamic, *A*daptive *S*treaming over *H*TTP

❖ *server*:

- divides video file into multiple chunks
- each chunk stored, encoded at different rates
- *manifest file*: provides URLs for different chunks


❖ *client*:

- periodically measures server-to-client bandwidth
- consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)

CDN content access: a closer look


Bob (client) requests video `http://netcinema.com/6Y7B23V`

- video stored in CDN at `http://KingCDN.com/NetC6y&B23V`



Chapter 3

Transport Layer



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

Slides adopted from original ones provided by the textbook authors.

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Transport vs. network layer

- ❖ **transport layer:** logical communication between processes
 - relies on, enhances, network layer services
 - **network layer:** logical communication between hosts
- ❖ two transport-layer protocols
 - TCP: reliable, in-order delivery
 - UDP: unreliable, unordered delivery

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Demultiplexing

- ❖ UDP socket identified by 2-tuple:
 - dest IP address
 - dest port number
- ❖ TCP socket identified by 4-tuple:
 - source IP address
 - source port number
 - dest IP address
 - dest port number

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management


3.6 principles of congestion control

3.7 TCP congestion control

UDP: User Datagram Protocol [RFC 768]

- ❖ “best effort” service, UDP segments may be:
 - lost
 - delivered out-of-order to app
- ❖ *connectionless*:
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others
- ❖ UDP use:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS
 - SNMP

UDP: segment header



UDP segment format

why is there a UDP?

- ❖ no connection establishment (which can add delay)
- ❖ simple: no connection state at sender, receiver
- ❖ small header size
- ❖ no congestion control: UDP can blast away as fast as desired

UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

sender:

- ❖ treat segment contents, including header fields, as sequence of 16-bit integers
- ❖ checksum: addition (one's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

receiver:

- ❖ compute checksum of received segment
- ❖ check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected.
But maybe errors nonetheless? More later
-

Internet checksum: example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	0	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control


Reliable Data Transfer

- ❖ **Challenge:** TCP requires reliable data transfer, but the underlying protocol IP is not reliable.
- ❖ Possible errors
 - channel with bit errors -> checksum, ACK/NAK, retransmission
 - with corrupted ACK/NAKs -> retransmission, sequence #
 - without NAKs -> ACK retransmission
 - channels with packet loss -> timer

Pipelined protocols

stop-and-wait: one packet at a time

pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts



(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

utilization: $U_{\text{sender}} = \frac{n L / R}{RTT + L / R}$

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control


3.7 TCP congestion control

TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

- ❖ **point-to-point:**
 - one sender, one receiver
- ❖ **reliable, in-order *byte stream*:**
 - no “message boundaries”
- ❖ **pipelined:**
 - TCP congestion and flow control set window size
- ❖ **full duplex data:**
 - bi-directional data flow in same connection
 - MSS: maximum segment size
- ❖ **connection-oriented:**
 - handshaking (exchange of control msgs) initializes sender, receiver state before data exchange
- ❖ **flow controlled:**
 - sender will not overwhelm receiver

TCP segment structure



Maximum segment size (MSS)

- ❖ MSS: maximum bytes of TCP payload
- ❖ Sequence #: byte-stream # of first byte in segment
- ❖ E.g. file size 500,000 bytes, MSS 1,000 bytes

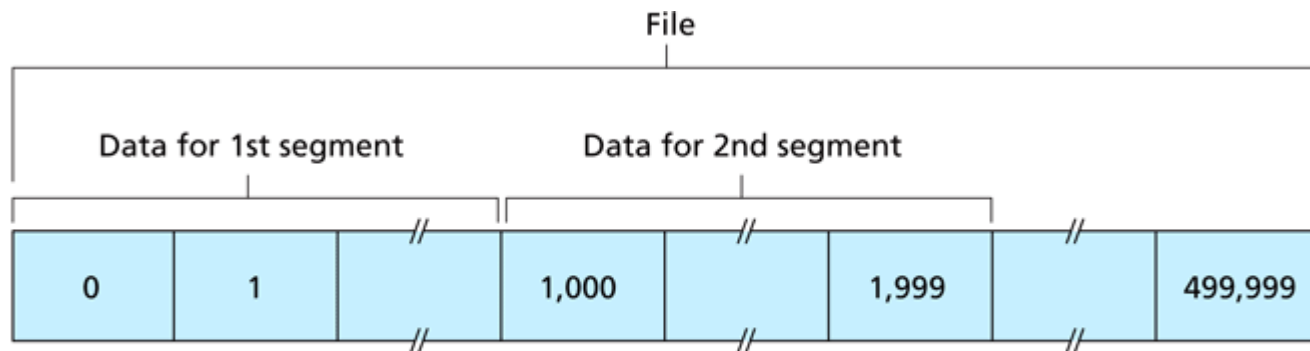


Figure 3.30 ♦ Dividing file data into TCP segments

TCP seq. #'s and ACKs

Seq. #'s:

- byte stream “number” of first byte in segment’s data

ACKs:

- seq # of next byte expected from other side
- cumulative ACK

Setting the time out

- $\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$
- $\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$
- $\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

TCP sender events:

data rcvd from app:

- ❖ create segment with seq #
- ❖ seq # is byte-stream number of first data byte in segment
- ❖ start timer if not already running
 - think of timer as for oldest unacked segment
 - expiration interval: `TimeoutInterval`

timeout:

- ❖ retransmit segment that caused timeout
- ❖ restart timer

ack rcvd:

- ❖ if ack acknowledges previously unacked segments
 - update ACK status
 - start timer if there are still unacked segments
 - triple duplicate ACKs: retransmit

TCP receiver events

<i>event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

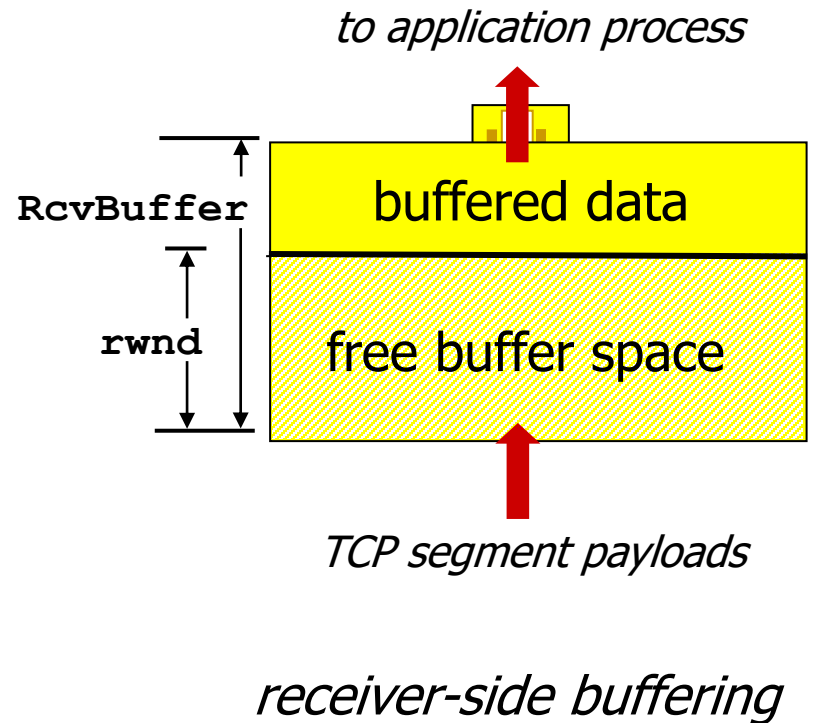
- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

TCP flow control

- ❖ receiver “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
 - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
 - many operating systems autoadjust **RcvBuffer**
- ❖ sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- ❖ guarantees receive buffer will not overflow



Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

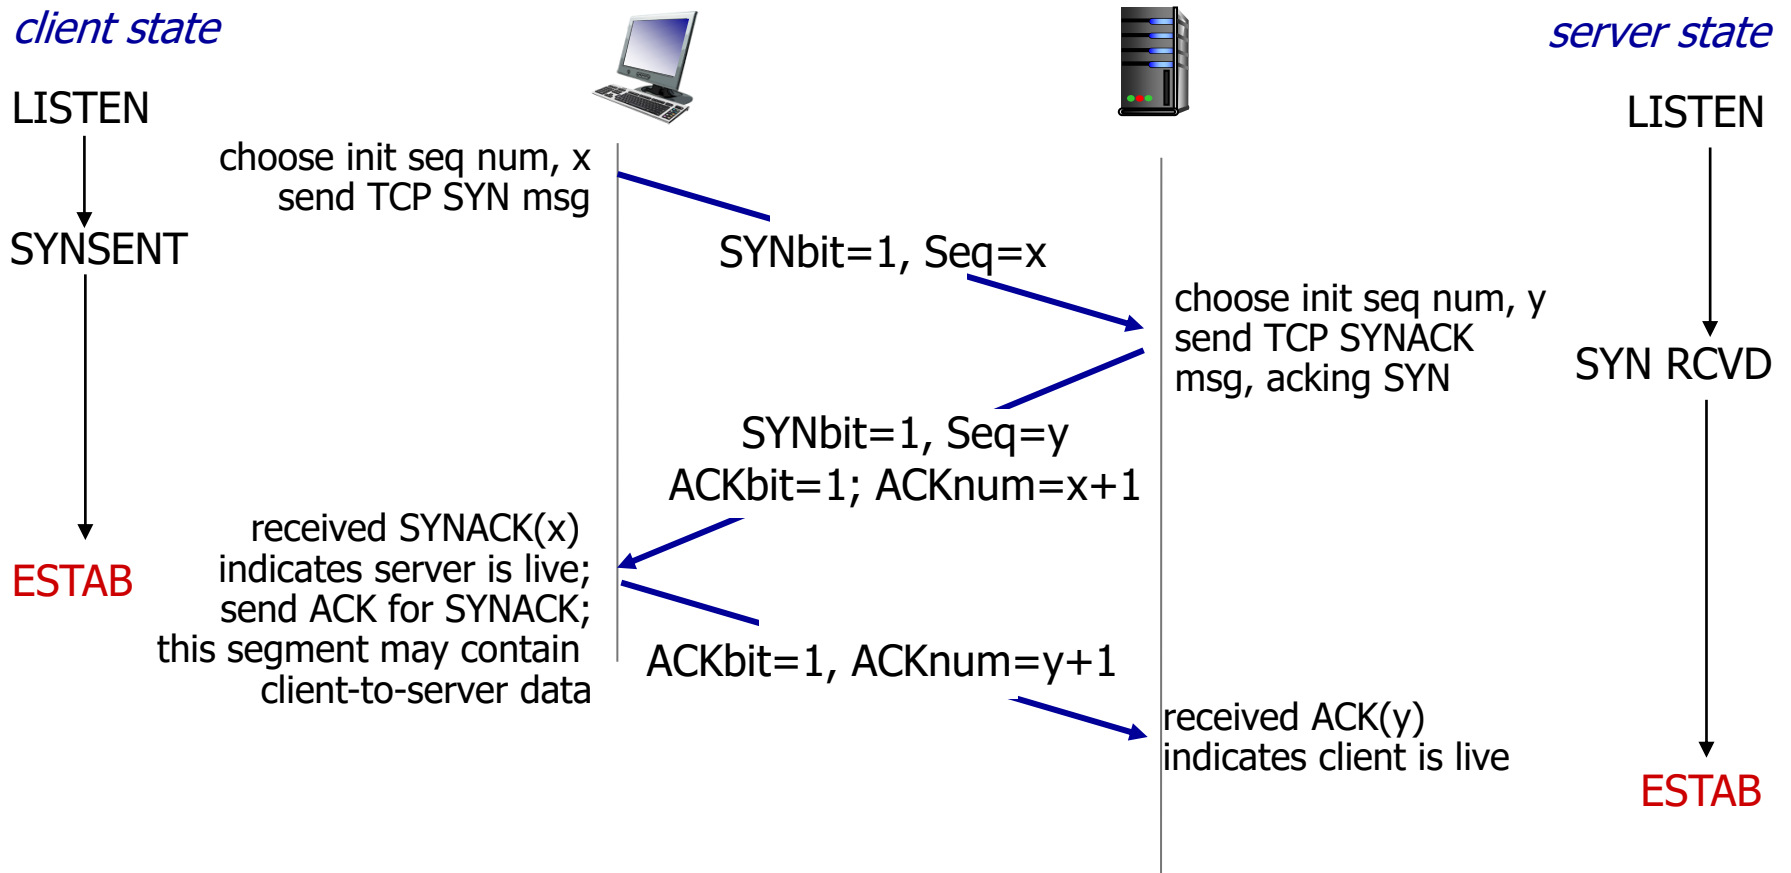
3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

TCP 3-way handshake



TCP: closing a connection

client state

ESTAB

`clientSocket.close()`

FIN_WAIT_1

can no longer
send but can
receive data

FIN_WAIT_2

wait for server
close

TIMED_WAIT

timed wait
for $2 * \text{max}$
segment lifetime

CLOSED



FINbit=1, seq=x

ACKbit=1; ACKnum=x+1

FINbit=1, seq=y

ACKbit=1; ACKnum=y+1

can still
send data

can no longer
send data

server state

ESTAB

CLOSE_WAIT

LAST_ACK

CLOSED

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Approaches towards congestion control

two broad approaches towards congestion control:

end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay
- ❖ approach taken by TCP

network-assisted congestion control:

- ❖ routers provide feedback to end systems
 - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
 - explicit rate for sender to send at

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Summary: TCP Congestion Control

- ❖ when `cwnd` < `ssthresh`, sender in **slow-start** phase, window grows exponentially.
- ❖ when `cwnd` ≥ `ssthresh`, sender is in **congestion-avoidance** phase, window grows linearly.
- ❖ when **triple duplicate ACK** occurs, `ssthresh` set to `cwnd/2`, `cwnd` set to `ssthresh + 3 MSS`
- ❖ when **timeout** occurs, `ssthresh` set to `cwnd/2`, `cwnd` set to 1 MSS.